


Implementing the Spatiocyte method on a Graphics Processing Unit (GPU)


Satya Arjunan

Laboratory for Biochemical Simulation
RIKEN Quantitative Biology Center

3 Sept 2016




Spatiocyte: a lattice-based particle simulator



Actual Spatiocyte simulation snapshots of kinesins (yellow and blue) and a microtubule

- Development started during my PhD in 2005 at Keio University
- Open source software available at <http://spatiocyte.org>
- Now part of E-Cell ver. 4 (Credits to Kazunari Kaizu and Suguru Kato)

Features of Spatiocyte



Arjunan & Tomita, Syst Synth Biol (2010)


- Spatiocyte can simulate diffusion and reaction of each molecule individually
- Reduces the computational time to resolve molecular collisions (reactions) by discretizing the simulation space into fine lattice voxels (size of a protein, 8 nm diameter)
- 1D, 2D and 3D stochastic reaction-diffusion

Examples of Spatiocyte application


E. coli MinE ring




Cell geometry directs PIP3/PTEN travelling waves in D. discoideum



Anionic lipid mediated spontaneous protein clustering and cooperative recruitment




Erythrocyte Band3 clustering



Arjunan et al., In preparation


Shimo et al., PLoS Comp Biol (2015)

Spontaneous polar localization of PAR-1/2 in C. elegans



Arjunan et al., In preparation

Microscopy bioimaging simulation




real simulated

Watabe et al., PLoS ONE (2015)

Arjunan et al., In progress

Computational issues of Spatiocyte



$$\langle r^2 \rangle = (2 r_v)^2$$

$$\tau_i = \frac{2r_v^2}{3D_i}$$

- Diffusion is the most costly operation
- Parallelize diffusion with GPUs
- Use C++ STL-like library of CUDA API called Thrust
- Code is available at <https://github.com/satya-arjunan/spatiocyte-cuda>



This repository Search

Pull requests Issues Gist



satya-arjunan / spatiocyte-cuda

Unwatch 1

Star 0

Fork 0

<> Code

Issues 0

Pull requests 0

Wiki

Pulse

Graphs

Settings

CUDA implementation of Spatiocyte — Edit

28 commits

1 branch

0 releases

1 contributor

Branch: master

New pull request

Create new file

Upload files

Find file


Clone or download

satya-arjunan Small change

Latest commit b41a6e3 4 days ago

Common.hpp	Removed boxes, simplified code	5 days ago
Compartment.cu	Moved device_vector offsets_ to Compartment	4 days ago
Compartment.hpp	Moved device_vector offsets_ to Compartment	4 days ago
Diffuser.cu	Small change	4 days ago
Diffuser.hpp	Moved device_vector offsets_ to Compartment	4 days ago
Lattice.cu	Moved lattice voxels into device_vector, faster initialization	4 days ago
Lattice.hpp	Moved lattice voxels into device_vector, faster initialization	4 days ago
Makefile	Update README and simplified Makefile	4 days ago
Model.cpp	Removed boxes, simplified code	5 days ago
Model.hpp	Removed boxes, simplified code	5 days ago
README.md	Update README and simplified Makefile	4 days ago
Spatocyte.cpp	Removed debug messages	4 days ago
Spatocyte.hpp	Sequential works, no bugs with large lattice	13 days ago

spatiocyte-cuda performance



1. Successfully implemented diffusion with GPU using CUDA Thrust on nVidia GTX 980
2. Avoided race conditions using atomic operation CAS (compare and swap)
3. Can simulate diffusion of 10,000,000 molecules with 5 nm diameter at $1\text{e-}12 \text{ m}^2\text{s}^{-1}$ for 10 s, which takes 11 hours of total run time
4. Achieved 125x speedup compared to serial implementation on Intel Xeon X5680 3.33 GHz

CUDA Toolkit v7.5

Thrust

1. Introduction

1.1. Installation and Versioning

2. Vectors

2.1. Thrust Namespace

2.2. Iterators and Static Dispatching

3. Algorithms

3.1. Transformations

3.2. Reductions

3.3. Prefix-Sums

3.4. Reordering

3.5. Sorting

4. Fancy Iterators

4.1. constant_iterator

4.2. counting_iterator

4.3. transform_iterator

4.4. permutation_iterator

4.5. zip_iterator

5. Additional Resources

Thrust ([PDF](#)) - v7.5 ([older](#)) - Last updated September 1, 2015 - [Send Feedback](#) - [f](#) [t](#) [in](#) [d](#) [t](#) [g](#)

Thrust

1. Introduction

Thrust is a C++ template library for CUDA based on the Standard Template Library (STL). Thrust allows you to implement high performance parallel applications with minimal programming effort through a high-level interface that is fully interoperable with CUDA C.

Thrust provides a rich collection of data parallel primitives such as scan, sort, and reduce, which can be composed together to implement complex algorithms with concise, readable source code. By describing your computation in terms of these high-level abstractions you provide Thrust with the freedom to select the most efficient implementation automatically. As a result, Thrust can be utilized in rapid prototyping of CUDA applications, where programmer productivity matters most, as well as in production, where robustness and absolute performance are crucial.

This document describes how to develop CUDA applications with Thrust. The tutorial is intended to be accessible, even if you have limited C++ or CUDA experience.

1.1. Installation and Versioning

Installing the CUDA Toolkit will copy Thrust header files to the standard CUDA include directory for your system. Since Thrust is a template library of header files, no further installation is necessary to start using Thrust.

In addition, new versions of Thrust continue to be available online through the GitHub [Thrust project page](#). The version of Thrust included in this version of the CUDA Toolkit corresponds to version 1.7.0 from the Thrust project page.

2. Vectors

Thrust provides two [vector](#) containers, [host_vector](#) and [device_vector](#). As the names suggest, [host_vector](#) is stored in host memory while [device_vector](#) lives in GPU device memory. Thrust's vector containers are just like [std::vector](#) in the [C++ STL](#). Like [std::vector](#), [host_vector](#) and [device_vector](#) are generic containers (able to store any data type) that can be resized dynamically. The following source code illustrates the use of Thrust's vector containers.

```
#include <thrust/host_vector.h>
#include <thrust/device_vector.h>

#include <iostream>

int main(void)
{
    // H has storage for 4 integers
    .....
```



```
void Diffuser::walk() {  
    const size_t size(mols_.size());  
    reacted_.resize(size);  
    thrust::transform(thrust::device,  
        thrust::counting_iterator<unsigned>(0),  
        thrust::counting_iterator<unsigned>(size),  
        mols_.begin(),  
        mols_.begin(),  
        generate(  
            seed_,  
            stride_,  
            id_stride_,  
            vac_id_,  
            thrust::raw_pointer_cast(&is_reactive_[0]),  
            thrust::raw_pointer_cast(&offsets_[0]),  
            thrust::raw_pointer_cast(&reacteds_[0]),  
            thrust::raw_pointer_cast(&voxels_[0])));  
    seed_ += size;  
}
```

```

struct generate {
    __host__ __device__ generate(
        const unsigned seed,
        const voxel_t stride,
        const voxel_t id_stride,
        const voxel_t vac_id,
        const bool* is_reactive,
        const mol_t* offsets,
        umol_t* reacteds,
        voxel_t* voxels):
    seed_(seed),
    stride_(stride),
    id_stride_(id_stride),
    vac_id_(vac_id),
    is_reactive_(is_reactive),
    offsets_(offsets),
    reacteds_(reacteds),
    voxels_(voxels) {}

    __device__ umol_t operator()(const unsigned index, const umol_t vdx) const {
        curandState s;
        curand_init(seed_+index, 0, 0, &s);
        float ranf(curand_uniform(&s)*11.999999);
        const unsigned rand((unsigned)truncf(ranf));
        const bool odd_lay((vdx/NUM_COLROW)&1);
        const bool odd_col((vdx%NUM_COLROW/NUM_ROW)&1);
        mol2_t val(mol2_t(vdx)+offsets_[rand+(24&(-odd_lay))+(12&(-odd_col))]);
        const voxel_t res(atomicCAS(voxels_+val, vac_id_, index+id_stride_));
        //If not occupied, walk:
        if(res == vac_id_) {
            voxels_[vdx] = vac_id_;
            reacteds_[index] = 0;
            return val;
        }
        return vdx;
    }
}

```